

Ahmed Sayed

Nov 27, 2017

08:38 PM -0600



Limit increase request 1
Service: EC2 Instances
Region: US West (Oregon)
Primary Instance Type: i3.large
Limit name: Instance Limit
New limit value: 9

Use case description: Hello, all

Currently, I am enrolled in a Cloud computing class which needs to perform Hadoop experiment on a cluster, So, I will be needing minimum 9 instances running at the same time.

Amazon Web Services

Nov 28, 2017

02:42 AM -0600



Hello Ahmed,

I trust you are having a lovely day! This is Rodney from AWS billing and accounts.

I'm currently working on your i3.large Instances limit increase request for 9 in the US West (Oregon) region. In this particular case, I have to collaborate with my Service Team to get approval.

I'm going to hold on to your case and the second I get an update from my Service Team I will let you know. I understand this is a really important resource for you so I will do my best to expedite this request.

Thank you very much for your patience while we work on this.

In the meantime, if there is anything you need, please feel free to reply to this case.

Have a lovely day!

Best regards,

Rodney W
Amazon Web Services

Check out the AWS Support Knowledge Center, a knowledge base of articles and videos that answer customer questions about AWS services:

https://aws.amazon.com/premiumsupport/knowledge-center/?icmpid=support_email_category

We value your feedback. Please rate my response using the link below.

Case Details

Basic Support Plan [Change](#)

Subject	Limit Increase: EC2 Instances
Case ID	4680901651
Created	Nov 27, 2017 08:38 PM -0600
Case type	Service Limits
By	asayed2@hawk.iit.edu
Status	Pending Customer Action
Severity	General guidance
Category	Service Limit Increase, EC2 Instances
CCd emails	

Correspondence

[Reply](#)

[Close Case](#)

Amazon Web Services

Nov 28, 2017
08:14 AM -0600



Hello,

Thank you for your patience while we reviewed your request.

I'm happy to inform you that we've approved and processed your i3.large instance limit increase request for the US West (Oregon) region, and your new limit is 9. Please keep in mind that it can sometimes take up to 15 minutes for the new limit to propagate and become available for use.

I hope this helps, but please let us know if you require further assistance.

Have a wonderful day!

Best regards,

Rodney W
Amazon Web Services

AWS Services Resource Groups Ahmed Sayed Oregon Support

1. Choose AMI 2. Choose Instance Type 3. Configure Instance 4. Add Storage 5. Add Tags 6. Configure Security Group 7. Review

Step 3: Configure Instance Details

Configure the instance to suit your requirements. You can launch multiple instances from the same AMI, request Spot instances to take advantage of the lower pricing, assign an access management role to the instance, and more.

Number of instances Launch into Auto Scaling Group

You may want to consider launching these instances into an Auto Scaling Group to help you maintain application availability and for easy scaling in the future. Learn how Auto Scaling can help your application stay healthy and cost effective.

Purchasing option Request Spot instances

Network Create new subnet

Subnet

Auto-assign Public IP

Placement group Add instance to placement group.

IAM role

Shutdown behavior

Enable termination protection Protect against accidental termination

Monitoring Enable CloudWatch detailed monitoring

Additional charges apply

[Cancel](#) [Previous](#) [Review and Launch](#) [Next: Add Storage](#)

The screenshot shows the AWS EC2 Dashboard. On the left, there's a sidebar with navigation links for EC2 Dashboard, Events, Tags, Reports, Limits, Instances (selected), Launch Templates, Spot Requests, Reserved Instances, Dedicated Hosts, Scheduled Instances, Images (AMIs, Bundle Tasks), and Elastic Block Store (Volumes, Snapshots). The main area has tabs for Launch Instance, Connect, and Actions. A search bar at the top says "Filter by tags and attributes or search by keyword". Below it is a table of instances:

Name	Instance ID	Instance Type	Availability Zone	Instance State	Status Checks	Alarm S	Public DNS (IPv4)
HadoopMaster	i-02034010c6e8fe8c4	i3.large	us-west-2c	running	2/2 checks ...	None	ec2-34-216-149-36.us-west-2.compute.amazonaws.com
HadoopSecondaryNameNo...	i-0375461866537a32	i3.large	us-west-2c	running	2/2 checks ...	None	ec2-34-216-149-189.us-west-2.compute.amazonaws.com
HadoopSlave1	i-04de512b6ca783a61	i3.large	us-west-2c	running	2/2 checks ...	None	ec2-34-208-66-218.us-west-2.compute.amazonaws.com
HadoopSlave2	i-05837693013e5d60f	i3.large	us-west-2c	running	2/2 checks ...	None	ec2-52-25-148-112.us-west-2.compute.amazonaws.com
HadoopSlave3	i-071efc8a0fc1b54a0	i3.large	us-west-2c	running	2/2 checks ...	None	ec2-35-163-241-78.us-west-2.compute.amazonaws.com
HadoopSlave4	i-0868816dcda67682d	i3.large	us-west-2c	running	2/2 checks ...	None	ec2-35-167-88-79.us-west-2.compute.amazonaws.com
HadoopSlave5	i-0bb89173b3243a91f	i3.large	us-west-2c	running	2/2 checks ...	None	ec2-34-216-136-103.us-west-2.compute.amazonaws.com
HadoopSlave6	i-0c324faddd891cc7c	i3.large	us-west-2c	running	2/2 checks ...	None	ec2-52-42-30-34.us-west-2.compute.amazonaws.com

Below the table, a specific instance is selected: **Instance: i-02034010c6e8fe8c4 (HadoopMaster)**. The Public DNS is **ec2-34-216-149-36.us-west-2.compute.amazonaws.com**. The detailed view shows the following information:

Description	Value
Instance ID	i-02034010c6e8fe8c4
Instance state	running
Instance type	i3.large
Elastic IPs	-
Availability zone	us-west-2c
Security groups	launch-wizard-29, view inbound rules
Scheduled events	No scheduled events
AMI ID	ubuntu/images/hvm-ssd/ubuntu-xenial-16.04-amd64-server-20171121.1 (ami-0def3275)
Public DNS (IPv4)	ec2-34-216-149-36.us-west-2.compute.amazonaws.com
IPv4 Public IP	34.216.149.36
Private IPs	-
Private DNS	ip-172-31-6-96.us-west-2.compute.internal
Secondary private IPs	172.31.6.96
VPC ID	vpc-61286606
Subnet ID	subnet-bf0414e

```
[ubuntu@ip-172-31-6-96:~$ ssh -i "AhmedShared.pem" ubuntu@ec2-34-216-149-189.us-west-2.compute.amazonaws.com
The authenticity of host 'ec2-34-216-149-189.us-west-2.compute.amazonaws.com' (172.31.6.255)
  can't be established.
ECDSA key fingerprint is SHA256:bJRNQSxgNWEzIxSFRw7kYuNkyIv1SrALv0+VqUiWThE.
[Are you sure you want to continue connecting (yes/no)? yes
Warning: Permanently added 'ec2-34-216-149-189.us-west-2.compute.amazonaws.com' (ECDSA) to the list of known hosts.
Welcome to Ubuntu 16.04.3 LTS (GNU/Linux 4.4.0-1041-aws x86_64)

 * Documentation:  https://help.ubuntu.com
 * Management:    https://landscape.canonical.com
 * Support:        https://ubuntu.com/advantage

 Get cloud support with Ubuntu Advantage Cloud Guest:
 http://www.ubuntu.com/business/services/cloud

5 packages can be updated.
0 updates are security updates.

Last login: Sun Dec  3 01:29:44 2017 from 207.237.204.157]
```

```

Are you sure you want to continue connecting (yes/no)? The authenticity of host 'ec2-52-25-148-112.us-west-2.compute.amazonaws.com (172.31.13.139)' can't be established.
ECDSA key fingerprint is SHA256:F6Bg0TAi6hr30mCUKRhSpaf0Kz+Sk0xVcnjjF/gopKI.
Are you sure you want to continue connecting (yes/no)? The authenticity of host 'ec2-34-208-66-218.us-west-2.compute.amazonaws.com (172.31.3.103)' can't be established.
ECDSA key fingerprint is SHA256:2GA3+JJjEzQdEqUhCy2RGvQTN7revbfDIv/iw5Nm3Js.
Are you sure you want to continue connecting (yes/no)? The authenticity of host 'ec2-52-42-30-34.us-west-2.compute.amazonaws.com (172.31.6.219)' can't be established.
ECDSA key fingerprint is SHA256:xdnj03GNTsHyhTvpklyF0e2Szdy83Ge+V6Mrp03bu08.
Are you sure you want to continue connecting (yes/no)? The authenticity of host 'ec2-35-167-88-79.us-west-2.compute.amazonaws.com (172.31.12.95)' can't be established.
ECDSA key fingerprint is SHA256:wyntTI9jSJpWAxCUTuS9RE+1SjuBvBKb63iCeqWp8Zs.
Are you sure you want to continue connecting (yes/no)? The authenticity of host 'ec2-34-216-136-103.us-west-2.compute.amazonaws.com (172.31.6.46)' can't be established.
ECDSA key fingerprint is SHA256:ulgULXFkxzKVNrzcEtCpB3FE5dvXSCFBfUljXyhb0.
Are you sure you want to continue connecting (yes/no)? ec2-34-216-149-189.us-west-2.compute.amazonaws.com: starting datanode, logging to /home/ubuntu/hadoop-2.7.4/logs/hadoop-ubuntu-datanode-ip-172-31-6-255.out

```

[yes]

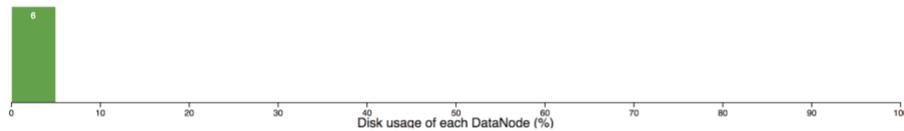
```

PA2 - ubuntu@ip-172-31-6-96: ~ ssh -i AhmedShared.pem ubuntu@ec2-34-215-219-2...
ec2-34-215-143-121.us-west-2.compute.amazonaws.com: starting nodemanage, logging to /home/ubuntu/hadoop-2.7.4/logs/yarn-ubuntu-nodemanager-ec2-34-215-143-121.us-west-2.compute.amazonaws.com.out
ec2-34-215-138-46.us-west-2.compute.amazonaws.com: starting nodemanage, logging to /home/ubuntu/hadoop-2.7.4/logs/yarn-ubuntu-nodemanager-ec2-34-215-138-46.us-west-2.compute.amazonaws.com.out
ec2-34-215-29-88.us-west-2.compute.amazonaws.com: starting nodemanage, logging to /home/ubuntu/hadoop-2.7.4/logs/yarn-ubuntu-nodemanager-ec2-34-215-29-88.us-west-2.compute.amazonaws.com.out
ec2-34-215-125-171.us-west-2.compute.amazonaws.com: starting nodemanage, logging to /home/ubuntu/hadoop-2.7.4/logs/yarn-ubuntu-nodemanager-ec2-34-215-125-171.us-west-2.compute.amazonaws.com.out
ec2-35-164-171-31.us-west-2.compute.amazonaws.com: starting nodemanage, logging to /home/ubuntu/hadoop-2.7.4/logs/yarn-ubuntu-nodemanager-ec2-35-164-171-31.us-west-2.compute.amazonaws.com.out
ec2-34-215-230-137.us-west-2.compute.amazonaws.com: starting nodemanage, logging to /home/ubuntu/hadoop-2.7.4/logs/yarn-ubuntu-nodemanager-ec2-34-215-230-137.us-west-2.compute.amazonaws.com.out
ubuntu@ip-172-31-6-96:~$ jps
26771 ResourceManager
26615 SecondaryNameNode
26392 NameNode
27038 Jps
ubuntu@ip-172-31-6-96:~$ 
PA3 - ubuntu@ip-172-31-13-139: ~ ssh -i AhmedShared.pem ubuntu@ec2-34-215-154-1...
ubuntu@ip-172-31-13-139:~$ export HADOOP_COMMON_HOME=$HADOOP_INSTALL
ubuntu@ip-172-31-13-139:~$ export HADOOP_HDFS_HOME=$HADOOP_INSTALL
ubuntu@ip-172-31-13-139:~$ ls
AhmedShared.pem
Cluster1.sh
hadoop-hadoop-2.7.4.tar.gz jdk1.8.0_131
Cluster1.sh hadoop-hadoop-2.7.4 jdk jdk-8u131-linux-x64.tar.gz
Agent pid 23184
ubuntu@ip-172-31-13-139:~$ ssh-add AhmedShared.pem
Identity added: AhmedShared.pem (AhmedShared)
ubuntu@ip-172-31-13-139:~$ mkdir hdfs
ubuntu@ip-172-31-13-139:~$ vi /etc/hadoop/slaves
ubuntu@ip-172-31-13-139:~$ vi /etc/hadoop/hadoop-env.sh
ubuntu@ip-172-31-13-139:~$ vi /etc/hadoop/hadoop-env.sh
ubuntu@ip-172-31-13-139:~$ jps
23634 Jps
2364 DataNode
ubuntu@ip-172-31-13-139:~$ jps
24448 DataNode
24393 Jps
ubuntu@ip-172-31-13-139:~$ 
PA4 - ubuntu@ip-172-31-13-139: ~ ssh -i AhmedShared.pem ubuntu@ec2-34-216-149-1...
ubuntu@ip-172-31-13-139:~$ export HADOOP_COMMON_HOME=$HADOOP_INSTALL
ubuntu@ip-172-31-13-139:~$ export HADOOP_HDFS_HOME=$HADOOP_INSTALL
ubuntu@ip-172-31-13-139:~$ export YARN_HOME=$HADOOP_INSTALL
ubuntu@ip-172-31-13-139:~$ ls
AhmedShared.pem
Cluster1.sh
hadoop-hadoop-2.7.4.tar.gz jdk1.8.0_131
Cluster1.sh hadoop-hadoop-2.7.4 jdk jdk-8u131-linux-x64.tar.gz
Agent pid 23184
ubuntu@ip-172-31-13-139:~$ ssh-add AhmedShared.pem
Identity added: AhmedShared.pem (AhmedShared)
ubuntu@ip-172-31-13-139:~$ mkdir hdfs
ubuntu@ip-172-31-13-139:~$ vi /etc/hadoop/slaves
ubuntu@ip-172-31-13-139:~$ vi /etc/hadoop/hadoop-env.sh
ubuntu@ip-172-31-13-139:~$ vi /etc/hadoop/hadoop-env.sh
ubuntu@ip-172-31-13-139:~$ jps
23717 Jps
23484 DataNode
ubuntu@ip-172-31-14-184:~$ jps
24134 DataNode
24397 Jps
ubuntu@ip-172-31-14-184:~$ 
PA5 - ubuntu@ip-172-31-13-139: ~ ssh -i AhmedShared.pem ubuntu@ec2-34-216-149-1...
ubuntu@ip-172-31-13-46:~$ export HADOOP_INSTALL=/home/ubuntu/hadoop
ubuntu@ip-172-31-13-46:~$ export PATH=$PATH:$HADOOP_INSTALL/bin
ubuntu@ip-172-31-13-46:~$ export HADOOP_MAPRED_HOME=$HADOOP_INSTALL
ubuntu@ip-172-31-13-46:~$ export HADOOP_COMMON_HOME=$HADOOP_INSTALL
ubuntu@ip-172-31-13-46:~$ export HADOOP_HDFS_HOME=$HADOOP_INSTALL
ubuntu@ip-172-31-13-46:~$ export YARN_HOME=$HADOOP_INSTALL
ubuntu@ip-172-31-13-46:~$ ls
AhmedShared.pem
Cluster1.sh
hadoop-hadoop-2.7.4.tar.gz jdk1.8.0_131
Cluster1.sh hadoop-hadoop-2.7.4 jdk jdk-8u131-linux-x64.tar.gz
Agent pid 23868
ubuntu@ip-172-31-13-46:~$ ssh-add AhmedShared.pem
Identity added: AhmedShared.pem (AhmedShared)
ubuntu@ip-172-31-13-46:~$ mkdir hdfs
ubuntu@ip-172-31-13-46:~$ vi /etc/hadoop/slaves
ubuntu@ip-172-31-13-46:~$ vi /etc/hadoop/hadoop-env.sh
ubuntu@ip-172-31-13-46:~$ jps
23687 Jps
23373 DataNode
ubuntu@ip-172-31-13-46:~$ jps
24276 Jps
24013 DataNode
ubuntu@ip-172-31-13-46:~$ 

```

Datanode Information

Datanode usage histogram



In operation

Node	Last contact	Admin State	Capacity	Used	Non DFS Used	Remaining	Blocks	Block pool used	Failed Volumes	Version
ec2-34-216-138-46.us-west-2.compute.amazonaws.com:50010 (172.31.3.103:50010)	1	In Service	290.76 GB	24 KB	2.78 GB	287.97 GB	0	24 KB (0%)	0	2.7.4
ec2-34-216-143-121.us-west-2.compute.amazonaws.com:50010 (172.31.12.95:50010)	1	In Service	290.76 GB	24 KB	2.78 GB	287.97 GB	0	24 KB (0%)	0	2.7.4
ec2-34-216-117-68.us-west-2.compute.amazonaws.com:50010 (172.31.6.255:50010)	1	In Service	290.76 GB	24 KB	2.78 GB	287.97 GB	0	24 KB (0%)	0	2.7.4
ec2-34-216-125-157.us-west-2.compute.amazonaws.com:50010 (172.31.14.184:50010)	1	In Service	290.76 GB	24 KB	2.78 GB	287.97 GB	0	24 KB (0%)	0	2.7.4
ec2-52-41-29-80.us-west-2.compute.amazonaws.com:50010 (172.31.6.46:50010)	1	In Service	290.76 GB	24 KB	2.78 GB	287.97 GB	0	24 KB (0%)	0	2.7.4
ec2-34-215-230-137.us-west-2.compute.amazonaws.com:50010 (172.31.6.219:50010)	1	In Service	290.76 GB	24 KB	2.78 GB	287.97 GB	0	24 KB (0%)	0	2.7.4

Decommissioning

Node	Last contact	Under replicated blocks	Blocks with no live replicas	Under Replicated Blocks In files under construction
------	--------------	-------------------------	------------------------------	--

Hadoop, 2017.

Hadoop	Overview	Datanodes	Datanode Volume Failures	Snapshot	Startup Progress	Utilities
--------	----------	-----------	--------------------------	----------	------------------	-----------

Overview 'ec2-34-215-219-28.us-west-2.compute.amazonaws.com:8020' (active)

Started:	Sat Dec 02 06:16:35 UTC 2017
Version:	2.7.4, rcd915e1e8d9d0131462a0b7301586c175728a282
Compiled:	2017-08-01T00:29Z by kshvachik from branch-2.7.4
Cluster ID:	CID-e605589a-f44d-4f34-bec8-c3e808d43ece
Block Pool ID:	BP-854754809-172.31.6.96-1512194588408

Summary

Security is off.
 Safemode is off.
 1 files and directories, 0 blocks = 1 total filesystem object(s).
 Heap Memory used 80.32 MB of 217.5 MB Heap Memory. Max Heap Memory is 889 MB.
 Non Heap Memory used 40.1 MB of 40.88 MB Committed Non Heap Memory. Max Non Heap Memory is -1 B.

Configured Capacity:	1.7 TB
DFS Used:	144 KB (0%)
Non DFS Used:	16.67 GB
DFS Remaining:	1.69 TB (99.04%)
Block Pool Used:	144 KB (0%)
DataNodes usages% (Min/Median/Max/stdDev):	0.00% / 0.00% / 0.00% / 0.00%
Live Nodes	6 (Decommissioned: 0)
Dead Nodes	0 (Decommissioned: 0)
Decommissioning Nodes	0
Total Datanode Volume Failures	0 (0 B)
Number of Under-Replicated Blocks	0
Number of Blocks Pending Deletion	0
Block Deletion Start Time	02/12/2017, 00:16:35

NameNode Journal Status

Current transaction ID: 6	
Journal Manager	State
FileJournalManager(root=/home/ubuntu/hdfstmp/dfs/name)	EditLogFileOutputStream(/home/ubuntu/hdfstmp/dfs/name/current/edits_inprogress_00000000000000000006)

NameNode Storage

Storage Directory	Type	State
/home/ubuntu/hdfstmp/dfs/name	IMAGE_AND_EDITS	Active

Hadoop, 2017.

```
[ubuntu@ip-172-31-6-96:~$ hadoop jar hadoop-2.7.4/share/hadoop/mapreduce/hadoop-mapreduce-examples-2.7.4.jar teragen 85800345 /teraInput8GB
17/12/03 01:41:48 INFO client.RMProxy: Connecting to ResourceManager at ec2-34-216-149-36.us-west-2.compute.amazonaws.com/172.31.6.96:9076
17/12/03 01:41:49 INFO terasort.TeraSort: Generating 85800345 using 2
17/12/03 01:41:50 INFO mapreduce.JobSubmitter: number of splits:2
17/12/03 01:41:50 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1512265063012_0001
17/12/03 01:41:50 INFO impl.YarnClientImpl: Submitted application application_1512265063012_0001
17/12/03 01:41:51 INFO mapreduce.Job: The url to track the job: http://ec2-34-216-149-36.us-west-2.compute.amazonaws.com:9078/proxy/application_1512265063012_0001/
17/12/03 01:41:51 INFO mapreduce.Job: Running job: job_1512265063012_0001
```

```
PA2 — ubuntu@ec2-34-215-219-28: ~ — ssh -i AhmedShared.pem ubuntu@ec2-34-215-2...
s-west-2.compute.amazonaws.com/172.31.6.96:9076
17/12/02 07:36:25 INFO impl.YarnClientImpl: Killed application application_1512196130380_0001
Killed job job_1512196130380_0001
[ubuntu@ip-172-31-6-96:~$ hadoop job -list
DEPRECATED: Use of this script to execute mapred command is deprecated.
Instead use the mapred command for it.

17/12/02 07:36:30 INFO client.RMProxy: Connecting to ResourceManager at ec2-34-215-219-28.us-west-2.compute.amazonaws.com/172.31.6.96:9076
Total jobs:0
      JobId      State          StartTime      UserName      Queue      P
 priority  UsedContainers  RsvdContainers  UsedMem      RsvdMem      NeededMem
 AM info
```

The screenshot shows the Hadoop Application Overview page. At the top, there's a navigation bar with links like 'Cluster', 'Nodes', 'Node Labels', 'Applications', 'Scheduler', 'Tools', and 'Help'. On the left, a sidebar lists cluster status: About, Nodes (NEW, NEW_SAVING, SUBMITTED, ACCEPTED, RUNNING, FINISHED, FAILED, KILLED), and Scheduler. The main content area is titled 'Application application_1512196130380_0001'. It displays application details: User: ubuntu, Name: TeraGen, Application Type: MAPREDUCE, Application Tags: (empty), YarnApplicationState: ACCEPTED: waiting for AM container to be allocated, launched and register with RM, Queue: default. It also shows the final status reported by AM: Application has not completed yet, Started: Sat Dec 02 06:41:38 +0000 2017, Elapsed: 3mins, 10sec, Tracking URL: ApplicationMaster, and Diagnostics. Below this is the 'Application Metrics' section, which includes tables for Total Resource Preempted, Total Number of Non-AM Containers Preempted, Total Number of AM Containers Preempted, Resource Preempted from Current Attempt, Number of Non-AM Containers Preempted from Current Attempt, and Aggregate Resource Allocation. At the bottom, there's a table showing application logs and a search bar.

Spark

10:40 AM -0400

Create Case

Case History

We have approved and processed your limit increase request(s). It can sometimes take up to 30 minutes for this to propagate and become available for use. I hope this helps, but please reopen this case if you encounter any issues.

Summary of limit(s) requested for increase:

[US East (Northern Virginia)]: EC2 Instances / Instance Limit (i3.4xlarge), New Limit = 2

Best regards,
Amazon Web Services

We value your feedback. Please rate our response below:

Was this response helpful? Click here to rate:
 Excellent

Feedback English (US)

© 2008 - 2017, Amazon Web Services, Inc. or its affiliates. All rights reserved. Privacy Policy Terms of Use

EC2 Management Console - Google Chrome

Secure | https://console.aws.amazon.com/ec2/v2/home?region=us-east-1#LaunchInstanceWizard:

aws Services Resource Groups

1. Choose AMI 2. Choose Instance Type 3. Configure Instance 4. Add Storage 5. Add Tags 6. Configure Security Group 7. Review

Step 2: Choose an Instance Type

Storage optimized	d2.8xlarge	36	244	24 x 2048	Yes	10 Gigabit	Yes
Storage optimized	i2.xlarge	4	30.5	1 x 800 (SSD)	Yes	Moderate	Yes
Storage optimized	i2.2xlarge	8	61	2 x 800 (SSD)	Yes	High	Yes
Storage optimized	i2.4xlarge	16	122	4 x 800 (SSD)	Yes	High	Yes
Storage optimized	i2.8xlarge	32	244	8 x 800 (SSD)	-	10 Gigabit	Yes
<input checked="" type="checkbox"/> Storage optimized	i3.large	2	15.25	1 x 475 (SSD)	Yes	Up to 10 Gigabit	Yes
Storage optimized	i3.xlarge	4	30.5	1 x 950 (SSD)	Yes	Up to 10 Gigabit	Yes
Storage optimized	i3.2xlarge	8	61	1 x 1900 (SSD)	Yes	Up to 10 Gigabit	Yes
Storage optimized	i3.4xlarge	16	122	2 x 1900 (SSD)	Yes	Up to 10 Gigabit	Yes
Storage optimized	i3.8xlarge	32	244	4 x 1900 (SSD)	Yes	10 Gigabit	Yes
Storage optimized	i3.16xlarge	64	488	8 x 1900 (SSD)	Yes	25 Gigabit	Yes

Cancel Previous Review and Launch Next: Configure Instance Details

Feedback English (US)

© 2008 - 2017, Amazon Web Services, Inc. or its affiliates. All rights reserved. Privacy Policy Terms of Use

IAM Management Console - Google Chrome

Inbox (1,062) - ssh... CS_553_Report - Go... IAM Management

Secure | https://console.aws.amazon.com/iam/home?region=us-east-1#groups

aWS Services Resource Groups

Hamdan Sher Global Support

Search IAM Create New Group Group Actions

Filter Showing 1 results

Group Name	Users	Inline Policy	Creation Time
SparkGroup	1		2017-11-22 18:02 CST

Dashboard Groups Users Roles Policies Identity providers Account settings Credential report Encryption keys

Feedback English (US)

© 2008 - 2017, Amazon Web Services, Inc. or its affiliates. All rights reserved. Privacy Policy Terms of Use

Detailed description: This screenshot shows the AWS IAM Management Console. The left sidebar has 'Groups' selected. The main area displays a table with one row for 'SparkGroup'. The table columns are 'Group Name', 'Users', 'Inline Policy', and 'Creation Time'. The 'Group Name' column shows 'SparkGroup', 'Users' shows '1', 'Inline Policy' is empty, and 'Creation Time' shows '2017-11-22 18:02 CST'.

EC2 Management Console - Google Chrome

Inbox (1,062) - ssh... CS_553_Report - Go... EC2 Management

Secure | https://console.aws.amazon.com/ec2/v2/home?region=us-east-1#KeyPairs:sort=keyName

aWS Services Resource Groups

Hamdan Sher N. Virginia Support

Instances Images Elastic Block Store Network & Security Key Pairs Load Balancing

Create Key Pair Import Key Pair Delete

Filter by attributes or search by keyword

Key pair name	Fingerprint
sher	46:15:46:b3:e4:1f:7b:6b:60:67:8e:bc:6c:da:7d:6f:50:f7:98:80
spark	9a:92:8f:1b:76:85:98:df:de:99:c6:cc:3b:14:be:86:f1:c4:73:a8

Select a key pair

Feedback English (US)

© 2008 - 2017, Amazon Web Services, Inc. or its affiliates. All rights reserved. Privacy Policy Terms of Use

Detailed description: This screenshot shows the AWS EC2 Management Console. The left sidebar has 'Key Pairs' selected. The main area displays a table with two rows for 'sher' and 'spark'. The table columns are 'Key pair name' and 'Fingerprint'. The 'Key pair name' column shows 'sher' and 'spark', and the 'Fingerprint' column shows their respective hex values.

EC2 Management Console - Google Chrome

Inbox (1,062) - ssh... CS_553_Report - Go EC2 Management C

Secure | https://console.aws.amazon.com/ec2/v2/home?region=us-east-1#SecurityGroups:sort=groupid

AWS Services Resource Groups

Hamdan Sher N. Virginia Support

EC2 Dashboard Events Tags Reports Limits

INSTANCES Instances Launch Templates Spot Requests Reserved Instances Dedicated Hosts Scheduled Instances

IMAGES AMIs Bundle Tasks

ELASTIC BLOCK STORE Volumes Snapshots

NETWORK & SECURITY Security Groups Elastic IPs

Create Security Group Actions

Filter by tags and attributes or search by keyword

1 to 10 of 10

Name	Group ID	Group Name	VPC ID	Description
sg-1189556c	default	vpc-bb3cdbdd	default VPC security group	
sg-15f65a60	SparkTest-slaves	vpc-bb3cdbdd	Spark EC2 group	
sg-25db6450	hadoop	vpc-bb3cdbdd	hadoop-1 created 2017-11-25T16:50:08.699-06:00	
sg-31ff9844	launch-wizard-3	vpc-bb3cdbdd	launch-wizard-3 created 2017-12-01T17:37:53.8...	
sg-4cf4cf39	Spark1-slaves	vpc-bb3cdbdd	Spark EC2 group	
sg-84fa56f1	SparkTest-master	vpc-bb3cdbdd	Spark EC2 group	
sg-a878d3dd	Spark1-master	vpc-bb3cdbdd	Spark EC2 group	
sg-a93a85dc	launch-wizard-1	vpc-bb3cdbdd	launch-wizard-1 created 2017-11-25T19:13:23.6...	
sg-e36c0996	TestSprk	vpc-bb3cdbdd	TestSprk	
sg-fa7ce58f	launch-wizard-2	vpc-bb3cdbdd	launch-wizard-2 created 2017-11-30T19:51:13.2...	

Select a security group above

Feedback English (US)

© 2008 - 2017, Amazon Web Services, Inc. or its affiliates. All rights reserved. Privacy Policy Terms of Use

The screenshot shows the AWS EC2 Management Console interface. The left sidebar includes links for EC2 Dashboard, Events, Tags, Reports, Limits, INSTANCES, Instances, Launch Templates, Spot Requests, Reserved Instances, Dedicated Hosts, Scheduled Instances, IMAGES, AMIs, Bundle Tasks, ELASTIC BLOCK STORE, Volumes, Snapshots, and NETWORK & SECURITY. Under SECURITY GROUPS, there is a link to 'Create Security Group' and an 'Actions' dropdown. The main content area displays a table of security groups with columns: Name, Group ID, Group Name, VPC ID, and Description. A message at the bottom says 'Select a security group above'. At the bottom of the page are links for Feedback, English (US), Privacy Policy, and Terms of Use.

IAM Management Console - Google Chrome

Inbox (1,062) - ssh... CS_553_Report - Go IAM Management C

Secure | https://console.aws.amazon.com/iam/home?region=us-east-1#users

AWS Services Resource Groups

Hamdan Sher Global Support

Search IAM

Dashboard Groups

Users Add user Delete user

Find users by username or access key Showing 1 result

User name	Groups	Access key age	Password age	Last activity	MFA
ssher1	SparkGroup	10 days	None	Today	Not enabled

Encryption keys

Feedback English (US)

© 2008 - 2017, Amazon Web Services, Inc. or its affiliates. All rights reserved. Privacy Policy Terms of Use

The screenshot shows the AWS IAM Management Console interface. The left sidebar includes links for Search IAM, Dashboard, Groups, Users (which is selected), Roles, Policies, Identity providers, Account settings, Credential report, and Encryption keys. The main content area shows a table of users with columns: User name, Groups, Access key age, Password age, Last activity, and MFA. A search bar at the top allows finding users by username or access key. A message at the bottom says 'Showing 1 result'. At the bottom of the page are links for Feedback, English (US), Privacy Policy, and Terms of Use.

EC2 Management Console - Google Chrome

Inbox (1,000) HADOOP EC2 Manag... Billing Manag... EC2 Manag... Error: Bro... (1) What's New 1000 gigab... TeraSort_...

Name	Instance ID	Instance Type	Availability Zone	Instance State	Status Checks	Alarm Status	Public DNS (IPv4)
	i-0d34c5209c517ec94	i3.4xlarge	us-east-1b	running	2/2 checks ...	None	ec2-52-206-51-120.co...
	i-0c8d308fec1a59d74	i3.large	us-east-1b	running	2/2 checks ...	None	ec2-54-82-138-51.com...

Instance: i-0d34c5209c517ec94 Public DNS: ec2-52-206-51-120.compute-1.amazonaws.com

Description Status Checks Monitoring Tags

Instance ID: i-0d34c5209c517ec94
 Instance state: running
 Instance type: i3.4xlarge
 Elastic IPs:

Public DNS (IPv4): ec2-52-206-51-120.compute-1.amazonaws.com
 IPv4 Public IP: 52.206.51.120
 IPv6 IPs: -
 Private DNS: ip-172-31-25-87.ec2.internal

Feedback English (US) © 2008 - 2017, Amazon Web Services, Inc. or its affiliates. All rights reserved. Privacy Policy Terms of Use Show all

```
ubuntu@ip-172-31-25-221:~ gensort hadoop makefile SharedMemory.class spark-1.6.0-bin-hadoop2.6
Gensort input128.txt scala SharedMemory.java spark-1.6.0-bin-hadoop2.6.tgz
ubuntu@ip-172-31-25-221:~$ rm input2.txt
ubuntu@ip-172-31-25-221:~$ rm input.txt
ubuntu@ip-172-31-25-221:~$ vt scode.scala
ubuntu@ip-172-31-25-221:~$ ./spark-1.6.0-bin-hadoop2.6/bin/spark-shell
log4j:WARN No appenders could be found for logger (org.apache.hadoop.metrics2.lib.MutableMetricsFactory).
log4j:WARN Please initialize the log4j system properly.
log4j:WARN See http://logging.apache.org/log4j/1.2/faq.html#noconfig for more info.
Using Spark's repl log4j profile: org/apache/spark/log4j-defaults-repl.properties
To adjust logging level use sc.setLogLevel("INFO")
Welcome to

      \ _ / \
     /_\ - \_ / / \_ \
    / . \_, _/_ / \_ \ \
   /_ / \
Using Scala version 2.10.5 (OpenJDK 64-Bit Server VM, Java 1.8.0_151)
Type in expressions to have them evaluated.
Type :help for more information.
Spark context available as sc.
17/12/02 06:38:09 WARN Connection: BoneCP specified but not present in CLASSPATH (or one of dependencies)
17/12/02 06:38:09 WARN Connection: BoneCP specified but not present in CLASSPATH (or one of dependencies)
17/12/02 06:38:14 WARN ObjectStore: Version information not found in metastore. hive.metastore.schema.verification is not enabled so recording the schema version 1.2.0
17/12/02 06:38:14 WARN ObjectStore: Failed to get database default, returning NoSuchObjectException
17/12/02 06:38:16 WARN Connection: BoneCP specified but not present in CLASSPATH (or one of dependencies)
17/12/02 06:38:16 WARN Connection: BoneCP specified but not present in CLASSPATH (or one of dependencies)
17/12/02 06:38:20 WARN ObjectStore: Version information not found in metastore. hive.metastore.schema.verification is not enabled so recording the schema version 1.2.0
17/12/02 06:38:20 WARN ObjectStore: Failed to get database default, returning NoSuchObjectException
SQL context available as sqlContext.

scala> :load /home/ubuntu/scode.scala
Loading /home/ubuntu/scode.scala...
lines: org.apache.spark.rdd.RDD[String] = MapPartitionsRDD[1] at textFile at <console>:27
ti: Long = 27005460953065
sort: org.apache.spark.rdd.RDD[String] = MapPartitionsRDD[7] at map at <console>:29
duration: Double = 2588.623730743
2588.623730743
scala>
```

```

ubuntu@ip-172-31-25-87:~ 
└── / - \ - \ - / - / \ - / version 1.6.0

Using Scala version 2.10.5 (OpenJDK 64-Bit Server VM, Java 1.8.0_151)
Type :help for more information.
Spark context available as sc.
17/12/02 07:32:42 WARN General: Plugin (Bundle) "org.datanucleus.api.jdo" is already registered. Ensure you dont have multiple JAR versions of the same plugin in the classpath. The URL "file:/home/ubuntu/spark-1.6.0-bin-hadoop2.6/lib/datanucleus-api-jdo-3.2.6.jar" is already registered, and you are trying to register an identical plugin located at URL "file:/home/ubuntu/spark/lib/datanucleus-api-jdo-3.2.6.jar."
17/12/02 07:32:42 WARN General: Plugin (Bundle) "org.datanucleus" is already registered. Ensure you dont have multiple JAR versions of the same plugin in the classpath. The URL "file:/home/ubuntu/spark-1.6.0-bin-hadoop2.6/lib/datanucleus-core-3.2.10.jar" is already registered, and you are trying to register an identical plugin located at URL "file:/home/ubuntu/spark/lib/datanucleus-core-3.2.10.jar."
17/12/02 07:32:42 WARN General: Plugin (Bundle) "org.datanucleus.store.rdbms" is already registered. Ensure you dont have multiple JAR versions of the same plugin in the classpath. The URL "file:/home/ubuntu/spark/lib/datanucleus-rdbms-3.2.9.jar" is already registered, and you are trying to register an identical plugin located at URL "file:/home/ubuntu/spark-1.6.0-bin-hadoop2.6/lib/datanucleus-rdbms-3.2.9.jar."
17/12/02 07:32:42 WARN Connection: BoneCP specified but not present in CLASSPATH (or one of dependencies)
17/12/02 07:32:45 WARN ObjectStore: Version information not found in metastore. hive.metastore.schema.verification is not enabled so recording the schema version 1.2.0
17/12/02 07:32:45 WARN ObjectStore: Failed to get database default, returning NoSuchObjectException
17/12/02 07:32:46 WARN General: Plugin (Bundle) "org.datanucleus.store.rdbms" is already registered. Ensure you dont have multiple JAR versions of the same plugin in the classpath. The URL "file:/home/ubuntu/spark/lib/datanucleus-rdbms-3.2.9.jar" is already registered, and you are trying to register an identical plugin located at URL "file:/home/ubuntu/spark-1.6.0-bin-hadoop2.6/lib/datanucleus-rdbms-3.2.9.jar."
17/12/02 07:32:46 WARN General: Plugin (Bundle) "org.datanucleus.api.jdo" is already registered. Ensure you dont have multiple JAR versions of the same plugin in the classpath. The URL "file:/home/ubuntu/spark-1.6.0-bin-hadoop2.6/lib/datanucleus-api-jdo-3.2.6.jar" is already registered, and you are trying to register an identical plugin located at URL "file:/home/ubuntu/spark/lib/datanucleus-api-jdo-3.2.6.jar."
17/12/02 07:32:46 WARN General: Plugin (Bundle) "org.datanucleus" is already registered. Ensure you dont have multiple JAR versions of the same plugin in the classpath. The URL "file:/home/ubuntu/spark-1.6.0-bin-hadoop2.6/lib/datanucleus-core-3.2.10.jar" is already registered, and you are trying to register an identical plugin located at URL "file:/home/ubuntu/spark/lib/datanucleus-core-3.2.10.jar."
17/12/02 07:32:46 WARN Connection: BoneCP specified but not present in CLASSPATH (or one of dependencies)
17/12/02 07:32:46 WARN ObjectStore: Version information not found in metastore. hive.metastore.schema.verification is not enabled so recording the schema version 1.2.0
SQL context available as sqlContext.

scala> :load /home/ubuntu/scode.scala
Loading /home/ubuntu/scode.scala...
lines: org.apache.spark.rdd.RDD[String] = MapPartitionsRDD[1] at textFile at <console>:27
ti: Long = 28284321391908
sort: org.apache.spark.rdd.RDD[String] = MapPartitionsRDD[7] at map at <console>:29
duration: Double = 803.935325867
803.935325867
[Stage 1]>                                     (0 + 16) / 4096

```

```

lastwalker@Chelsea: ~/Documents/SEMESTER-3/Cloud_Computing/A2/SPARK/spark/ec2
Connecting to s3.amazonaws.com (s3.amazonaws.com)|52.216.130.237|:80... connected.
HTTP request sent, awaiting response... 200 OK
Length: 30531249 (29M) [application/x-compressed]
Saving to: "scala-2.10.3.tgz"

scalar-2.10.3.tgz          100%[=====] 29.12M 68.9MB/s in 0.4s
2017-12-02 02:09:45 (68.9 MB/s) - 'scala-2.10.3.tgz' saved [30531249/30531249]

[timing] scala init: 00h 00m 00s
Initializing spark
--2017-12-02 02:09:45-- http://s3.amazonaws.com/spark-related-packages/spark-1.6.2-bin-hadoop1.tgz
Resolving s3.amazonaws.com (s3.amazonaws.com)... 52.216.130.237
Connecting to s3.amazonaws.com (s3.amazonaws.com)|52.216.130.237|:80... connected.
HTTP request sent, awaiting response... 404 Not Found
2017-12-02 02:09:46 ERROR 404: Not Found.

ERROR: Unknown Spark version
spark/init.sh: line 137: return: -1: invalid option
return: usage: return [n]
Unpacking Spark
tar (child): spark-*.tgz: Cannot open: No such file or directory
tar (child): Error is not recoverable: exiting now
tar: Child returned status 2
tar: Error is not recoverable: exiting now
rm: cannot remove 'spark-*.tgz': No such file or directory
mv: missing destination file operand after 'spark'
Try 'mv --help' for more information.
[timing] spark init: 00h 00m 01s
Initializing ephemeral-hdfs
--2017-12-02 02:09:46-- http://s3.amazonaws.com/spark-related-packages/hadoop-1.0.4.tar.gz
Resolving s3.amazonaws.com (s3.amazonaws.com)... 52.216.130.237
Connecting to s3.amazonaws.com (s3.amazonaws.com)|52.216.130.237|:80... connected.
HTTP request sent, awaiting response... 200 OK
Length: 62793050 (60M) [application/x-gzip]
Saving to: 'hadoop-1.0.4.tar.gz'

hadoop-1.0.4.tar.gz          100%[=====] 59.88M 68.7MB/s in 0.9s
2017-12-02 02:09:46 (68.7 MB/s) - 'hadoop-1.0.4.tar.gz' saved [62793050/62793050]

Unpacking Hadoop
RSYNC'ing /root/ephemeral-hdfs to slaves...

```

```
lastwalker@Chelsea: ~/Documents/SEMESTER-3/Cloud_Computing/A2/SPARK/spark/ec2
Creating local config files...
Connection to ec2-34-235-126-195.compute-1.amazonaws.com closed.
Configuring /root/spark/conf/spark-defaults.conf
Configuring /root/spark/conf/core-site.xml
Configuring /root/spark/conf/spark-env.sh
Configuring /root/persistent-hdfs/conf/masters
Configuring /root/persistent-hdfs/conf/slaves
Configuring /root/persistent-hdfs/conf/mapred-site.xml
Configuring /root/persistent-hdfs/conf/hadoop-env.sh
Configuring /root/persistent-hdfs/conf/core-site.xml
Configuring /root/persistent-hdfs/conf/hdfs-site.xml
Configuring /root/nareduce/hadoop.version
Configuring /root/nareduce/conf/masters
Configuring /root/nareduce/conf/slaves
Configuring /root/nareduce/conf/mapred-site.xml
Configuring /root/nareduce/conf/hadoop-env.sh
Configuring /root/nareduce/conf/core-site.xml
Configuring /root/nareduce/conf/hdfs-site.xml
Configuring /root/ephemeral-hdfs/conf/yarn-site.xml
Configuring /root/ephemeral-hdfs/conf/capacity-scheduler.xml
Configuring /root/ephemeral-hdfs/conf/masters
Configuring /root/ephemeral-hdfs/conf/slaves
Configuring /root/ephemeral-hdfs/conf/mapred-site.xml
Configuring /root/ephemeral-hdfs/conf/hadoop-metrics2.properties
Configuring /root/ephemeral-hdfs/conf/hadoop-env.sh
Configuring /root/ephemeral-hdfs/conf/core-site.xml
Configuring /root/ephemeral-hdfs/conf/hdfs-site.xml
Configuring /root/ephemeral-hdfs/conf/yarn-env.sh
Configuring /root/tachyon/conf/slaves
Configuring /root/tachyon/conf/workers
Configuring /root/tachyon/conf/tachyon-env.sh
Configuring /etc/httpd/conf/httpd.conf
Configuring /etc/httpd/conf.d/ganglia.conf
Configuring /etc/ganglia/gmetad.conf
Configuring /etc/ganglia/gmond.conf
Deploying Spark config files...
RSYNC'ing /root/spark/conf to slaves...
ec2-34-235-126-195.compute-1.amazonaws.com
ec2-54-211-21-138.compute-1.amazonaws.com
ec2-54-208-186-165.compute-1.amazonaws.com
ec2-34-227-191-19.compute-1.amazonaws.com
ec2-34-235-126-184.compute-1.amazonaws.com
```

```
lastwalker@Chelsea: ~/Documents/SEMESTER-3/Cloud_Computing/A2/SPARK/spark/ec2
17/12/02 02:12:22 INFO util.GSet: VM type          = 64-bit
17/12/02 02:12:22 INFO util.GSet: 2% max memory = 17.78 MB
17/12/02 02:12:22 INFO util.GSet: Capacity        = 2^21 = 2097152 entries
17/12/02 02:12:22 INFO util.GSet: recommended=2097152, actual=2097152
17/12/02 02:12:22 INFO namenode.FSNamesystem: fsOwner=root
17/12/02 02:12:22 INFO namenode.FSNamesystem: supergroup=supergroup
17/12/02 02:12:22 INFO namenode.FSNamesystem: lsPermissionEnabled=false
17/12/02 02:12:22 INFO namenode.FSNamesystem: dfs.block.invalidate.limit=100
17/12/02 02:12:22 INFO namenode.FSNamesystem: lsAccessTokenEnabled=false accessKeyUpdateInterval=0 min(s), accessTokenLifetime=0 min(s)
17/12/02 02:12:22 INFO namenode.NameNode: Caching file names occurring more than 10 times
17/12/02 02:12:22 INFO common.Storage: Image file of size 110 saved in 0 seconds.
17/12/02 02:12:23 INFO common.Storage: Storage directory /vol/persistent-hdfs/dfs/name has been successfully formatted.
17/12/02 02:12:23 INFO namenode.NameNode: SHUTDOWN_MSG:
*****SHUTDOWN_MSG: Shutting down NameNode at ip-172-31-4-93.ec2.internal/172.31.4.93
*****SHUTDOWN_MSG: Shutting down NameNode at ip-172-31-4-93.ec2.internal/172.31.4.93
Persistent HDFS installed, won't start by default...
[timing] persistent-hdfs setup: 00h 00m 07s
Setting up spark-standalone
RSYNC'ing /root/spark/conf to slaves...
ec2-34-235-126-195.compute-1.amazonaws.com
ec2-54-211-21-138.compute-1.amazonaws.com
ec2-54-208-186-165.compute-1.amazonaws.com
ec2-34-227-191-19.compute-1.amazonaws.com
ec2-34-235-126-184.compute-1.amazonaws.com
ec2-34-224-222-198.compute-1.amazonaws.com
ec2-34-203-212-89.compute-1.amazonaws.com
RSYNC'ing /root/spark-ec2 to slaves...
ec2-34-235-126-195.compute-1.amazonaws.com
ec2-54-211-21-138.compute-1.amazonaws.com
ec2-54-208-186-165.compute-1.amazonaws.com
ec2-34-227-191-19.compute-1.amazonaws.com
ec2-34-235-126-184.compute-1.amazonaws.com
ec2-34-224-222-198.compute-1.amazonaws.com
ec2-34-203-212-89.compute-1.amazonaws.com
./spark-standalone/setup.sh: line 22: /root/spark/sbin/stop-all.sh: No such file or directory
./spark-standalone/setup.sh: line 27: /root/spark/sbin/start-master.sh: No such file or directory
[timing] spark-standalone setup: 00h 00m 36s
Setting up tachyon
RSYNC'ing /root/tachyon to slaves...
ec2-34-235-126-195.compute-1.amazonaws.com
ec2-54-211-21-138.compute-1.amazonaws.com
```

EC2 Management Console - Google Chrome

Inbox (1,000) HADOOP EC2 Manager Billing Manager EC2 Manager Error: Bro WhatsApp 64 gigabyte Amazon E 8:17 PM

Secure | https://console.aws.amazon.com/ec2/v2/home?region=us-east-1#

aws Services Resource Groups

EC2 Dashboard Resources Account Attributes

You are using the following Amazon EC2 resources in the US East (N. Virginia) region:

10 Running Instances	0 Elastic IPs
0 Dedicated Hosts	0 Snapshots
10 Volumes	0 Load Balancers
2 Key Pairs	10 Security Groups
0 Placement Groups	

EC2 Spot. Save up to 90% off On-Demand Prices. Turbo Boost your Workloads. Get started with Amazon EC2 Spot Instances.

Create Instance

To start using Amazon EC2 you will want to launch a virtual server, known as an Amazon EC2 instance.

Launch Instance

Note: Your instances will launch in the US East (N. Virginia) region

Service Health Scheduled Events

Service Status: US East (N. Virginia):

AWS Marketplace

Find free software trial products in the AWS Marketplace from the EC2 Launch Wizard. Or try these popular AMIs:

ParrotLabs NextGen Firewall F-Secure

Feedback English (US)

prog2_report....pdf Removed terasort_rep....pdf terasort_report.pdf Report (1).docx prog2_report.pdf Show all

2008 - 2017, Amazon Web Services, Inc. or its affiliates. All rights reserved. Privacy Policy Terms of Use

EC2 Management Console - Google Chrome

Inbox HADOOP Tera EC2 Billing EC2 Error (1) V (7) H 64 g Amazon scal EC2 10:21 PM

Secure | https://console.aws.amazon.com/ec2/v2/home?region=us-east-1#instances:sort=statusChecks

aws Services Resource Groups

EC2 Dashboard Instances Actions

Launch Instance Connect Actions

Filter by tags and attributes or search by keyword

Name	Instance ID	Instance Type	Availability Zone	Instance State	Status Checks	Alarm Status
SparkTest-slave-i-015aba238f99454ff	i-015aba238f99454ff	i3.large	us-east-1a	running	2/2 checks ...	None
SparkTest-master-i-02518eaa65142e...	i-02518eaa65142e6...	i3.large	us-east-1a	running	2/2 checks ...	None
SparkTest-slave-i-028dd445b3df69a1	i-028dd445b3df69a1	i3.large	us-east-1a	running	2/2 checks ...	None
SparkTest-slave-i-04c415f6b921a3b1	i-04c415f6b921a3b1	i3.large	us-east-1a	running	2/2 checks ...	None
SparkTest-slave-i-057cef2ad1420e845	i-057cef2ad1420e845	i3.large	us-east-1a	running	2/2 checks ...	None
SparkTest-slave-i-0ab67b36cd993e2aa	i-0ab67b36cd993e2aa	i3.large	us-east-1a	running	2/2 checks ...	None
	i-0c8d308fec1a59d74	i3.large	us-east-1b	running	2/2 checks ...	None
	i-0d34c5209c517ec94	i3.4xlarge	us-east-1b	running	2/2 checks ...	None
	SparkTest-slave-i-0dafef757901c334db	i3.large	us-east-1a	running	2/2 checks ...	None
	i-0e56886cff13d6ca	i3.large	us-east-1a	running	2/2 checks ...	None

Select an instance above

Feedback English (US)

prog2_report....pdf Removed terasort_rep....pdf terasort_report.pdf Report (1).docx prog2_report.pdf Show all

2008 - 2017, Amazon Web Services, Inc. or its affiliates. All rights reserved. Privacy Policy Terms of Use